



PANDAS DATA STRUCTURES



Data Structures in Pandas

- Pandas has two main data structures;
 - DataFrame, which is two dimensional
 - Series, which is one dimensional

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
5	Apple Cinnamon Cheerios	110	2	25	29.509541
6	Apple Jacks	110	2	25	33.174094
7	Basic 4	130	3	25	37.038562
8	Bran Chex	90	2	25	49.120253
9	Bran Flakes	90	3	25	53.313813

DataFrame

```
0    70
1   120
2    70
3    50
4   110
5   110
6   110
7   130
8    90
9    90
Name: calories, dtype: int64
```

Series



What is Pandas DataFrame

- Pandas provides a two dimensional data structure called DataFrame.
- A row is represented by row labels, also called index, which may be numerical or strings.
- A column is represented by column labels which may be numerical or strings.
- Following DataFrame contains 10 rows (0-9) and 5 columns (name, calories, protein, vitamins, rating)

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
5	Apple Cinnamon Cheerios	110	2	25	29.509541
6	Apple Jacks	110	2	25	33.174094
7	Basic 4	130	3	25	37.038562
8	Bran Chex	90	2	25	49.120253
9	Bran Flakes	90	3	25	53.313813

index

column labels



What is Pandas Series

- A Series in Pandas is a one dimensional data structure.
- It consists of a single row or column.
- Following Series contains 10 rows (0-9) and 1 column called calories.

0	70
1	120
2	70
3	50
4	110
5	110
6	110
7	130
8	90
9	90
Name: calories, dtype: int64	

index

column values

column name

column data type



DataFrame and Series

- A Pandas DataFrame is just a collection of one or more Series.
- The Series in the previous example was extracted from the DataFrame.

DataFrame →

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
5	Apple Cinnamon Cheerios	110	2	25	29.509541
6	Apple Jacks	110	2	25	33.174094
7	Basic 4	130	3	25	37.038562
8	Bran Chex	90	2	25	49.120253
9	Bran Flakes	90	3	25	53.313813

→ *Series*

```
0    70
1   120
2    70
3    50
4   110
5   110
6   110
7   130
8    90
9    90
Name: calories, dtype: int64
```

Identical



Creating a DataFrame Using Lists (1/3)

- We can create a DataFrame using lists.
- We pass the list as an argument to the `pandas.DataFrame()` function which returns us a DataFrame.
- Pandas automatically assigns numerical row labels to each row of the DataFrame.
- Since, we did not provide column labels, Pandas automatically assigned numerical column labels to each column as well.

```

In [2]: myList = [['Apple', 'Red'],
                  ['Banana', 'Yellow'],
                  ['Orange', 'Orange']]

In [3]: myDataFrame = pd.DataFrame(myList)
myDataFrame

Out[3]:
```

	0	1
0	Apple	Red
1	Banana	Yellow
2	Orange	Orange



Creating a DataFrame Using Lists (2/3)

- Let's create another DataFrame using the same list, but this time with custom column labels.
- `Pandas.DataFrame()` takes another optional argument called 'columns' which takes a list of custom column names to be set as columns' labels.

```
In [2]: myList = [['Apple', 'Red'],  
                  ['Banana', 'Yellow'],  
                  ['Orange', 'Orange']]
```

```
In [4]: myDataFrame = pd.DataFrame(myList, columns=['Fruit', 'Color'])  
myDataFrame
```

Out[4]:

	Fruit	Color
0	Apple	Red
1	Banana	Yellow
2	Orange	Orange



Creating a DataFrame Using Lists (3/3)

- As we know that a NumPy Array is similar to a Python list with added functionality, we can also convert a NumPy Array to a DataFrame using the same method.

```
In [9]: myList = np.array([[0, 1],  
                           [2, 3],  
                           [4, 5]])
```

```
In [10]: myDataFrame = pd.DataFrame(myList, columns=['even', 'odd'])  
myDataFrame
```

Out[10]:

	even	odd
0	0	1
1	2	3
2	4	5



Creating a DataFrame Using Dictionary

- We can also pass a dictionary to the `pandas.DataFrame()` function to create a DataFrame.
- Each key of the array should have a list of one or more values associated with it.
- The keys of the dictionary become column labels.
- Pandas automatically assigns numerical row labels to each row of the DataFrame.

The image shows a Jupyter Notebook interface with the title "DataFrames". The top bar includes a "Logout" button and a "Checkpoint created: 19:54:28" status. The menu bar contains "File", "Edit", "View", "Insert", "Cell", "Kernel", "Widgets", and "Help". Below the menu is a toolbar with icons for file operations and execution. The code area contains three input cells:

```
In [10]: import pandas as pd
import numpy as np

In [14]: myDictionary = {'Fruit': ['Apple', 'Banana', 'Orange'], 'Color': ['Red', 'Yellow', 'Orange']}
```

The third cell is selected, showing its output:

```
In [15]: myDataFrame = pd.DataFrame(myDictionary)
myDataFrame

Out[15]:
```

	Fruit	Color
0	Apple	Red
1	Banana	Yellow
2	Orange	Orange



Loading csv File as a DataFrame

- We can also load a csv (comma separated values) file as a DataFrame in Pandas using the `pandas.read_csv()` function.
- Each value of the first row of the csv file becomes a column label.
- Pandas automatically assigns numerical row labels to each row of the DataFrame.

```
In [11]: df = pd.read_csv('cereals.csv')  
df
```

```
Out[11]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
5	Apple Cinnamon Cheerios	110	2	25	29.509541
6	Apple Jacks	110	2	25	33.174094
7	Basic 4	130	3	25	37.038562
8	Bran Chex	90	2	25	49.120253
9	Bran Flakes	90	3	25	53.313813



Changing the Index Column

- We can set one of the existing columns as the new index column of the DataFrame using `.set_index()` function.

```
In [13]: df
```

```
Out[13]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843

```
In [14]: df.set_index('name')
```

```
Out[14]:
```

	calories	protein	vitamins	rating
name				
100% Bran	70	4	25	68.402973
100% Natural Bran	120	3	0	33.983679
All-Bran	70	4	25	59.425505
All-Bran with Extra Fiber	50	4	25	93.704912
Almond Delight	110	2	25	34.384843



Inplace (1/2)

- Remember that most of the functions in Pandas do not change the original DataFrame.
- In the previous section we changed the index column of our DataFrame. If we print our DataFrame again, we see that the original DataFrame is unchanged.

```
In [15]: df
```

```
Out[15]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843



Inplace (2/2)

- We can use the inplace argument to make changes to the original DataFrame.
- In the following example we use the .set_index() function to change the index of our DataFrame, and set inplace = True.
- As shown in the figure, our original DataFrame has been changed.

```
In [16]: df.set_index('name', inplace=True)
```

```
In [17]: df
```

```
Out[17]:
```

	calories	protein	vitamins	rating
name				
100% Bran	70	4	25	68.402973
100% Natural Bran	120	3	0	33.983679
All-Bran	70	4	25	59.425505
All-Bran with Extra Fiber	50	4	25	93.704912
Almond Delight	110	2	25	34.384843



Examining the Data (1/2)

head()

- head() function gives us the **first** 5 rows of the DataFrame/Series by default.
- To get more rows, we can pass the desired number as an argument to the head() function.

```
In [20]: df.head(7)
```

```
Out[20]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
5	Apple Cinnamon Cheerios	110	2	25	29.509541
6	Apple Jacks	110	2	25	33.174094



Examining the Data (2/2)

`tail()`

- `tail()` function gives us the **last** 5 rows of the DataFrame/Series by default.
- To get more rows, we can pass the desired number as an argument to the `tail()` function.

```
In [22]: df.tail(7)
```

```
Out[22]:
```

	name	calories	protein	vitamins	rating
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
5	Apple Cinnamon Cheerios	110	2	25	29.509541
6	Apple Jacks	110	2	25	33.174094
7	Basic 4	130	3	25	37.038562
8	Bran Chex	90	2	25	49.120253
9	Bran Flakes	90	3	25	53.313813



Statistical Summary

- We can use the describe() function to get a quick statistical summary of each column of the DataFrame.

```
In [25]: df.describe()
```

```
Out[25]:
```

	calories	protein	vitamins	rating
count	5.000000	5.000000	5.00000	5.000000
mean	84.000000	3.400000	20.00000	57.980382
std	29.664794	0.894427	11.18034	25.097570
min	50.000000	2.000000	0.00000	33.983679
25%	70.000000	3.000000	25.00000	34.384843
50%	70.000000	4.000000	25.00000	59.425505
75%	110.000000	4.000000	25.00000	68.402973
max	120.000000	4.000000	25.00000	93.704912



[] Operator for Row Slicing (1/2)

- We can use the brackets ([]) operator to slice rows of the DataFrame.
- We pass a start index (inclusive) and an end index (exclusive) to the bracket operator ([]) to slice the rows of the DataFrame.

```
In [26]: df[1:4]
```

```
Out[26]:
```

	name	calories	protein	vitamins	rating
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912



[] Operator for Row Slicing (2/2)

- Remember that [] operator works on row position and not row labels.
- For example, in the following case row labels are strings. But we pass positions of the rows that we want to slice.

```
In [30]: df
```

```
Out[30]:
```

	calories	protein	vitamins	rating
name				
100% Bran	70	4	25	68.402973
100% Natural Bran	120	3	0	33.983679
All-Bran	70	4	25	59.425505
All-Bran with Extra Fiber	50	4	25	93.704912
Almond Delight	110	2	25	34.384843

```
In [31]: df[1:4]
```

```
Out[31]:
```

	calories	protein	vitamins	rating
name				
100% Natural Bran	120	3	0	33.983679
All-Bran	70	4	25	59.425505
All-Bran with Extra Fiber	50	4	25	93.704912



Quiz Time

1. Consider the given DataFrame called df. Which of the following will give us rows 5-10 of the DataFrame?
 - a) `df[5:10]`
 - b) `df[5:11]`
 - c) `df[4:10]`
 - d) `df[4:11]`



Quiz Time

1. Consider the given DataFrame called df. Which of the following will give us rows 5-10 of the DataFrame?
 - a) df[5:10]
 - b) df[5:11]
 - c) df[4:10]
 - d) df[4:11]



[] Operator for Column Indexing

- We can also use the brackets ([]) operator to index column of the DataFrame.
- Indexing a single column returns a Series.
- Indexing a list of columns returns a DataFrame.
- Remember that for indexing columns, we pass their labels to the [] operator and not their positions.

```
In [34]: df[['name', 'rating']]
```

```
Out[34]:
```

	name	rating
0	100% Bran	68.402973
1	100% Natural Bran	33.983679
2	All-Bran	59.425505
3	All-Bran with Extra Fiber	93.704912
4	Almond Delight	34.384843



Boolean List

- We can also pass a list of booleans to the [] operator.
- We get all the rows of the DataFrame for which the corresponding element in the list is True.
- Rows of the DataFrame for which the corresponding element in the list is False are ignored.
- Note: Original DataFrame remains unchanged.

```
In [71]: df
```

```
Out[71]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843

```
In [75]: thirdRow = [False, False, True, False, False]  
df[thirdRow]
```

```
Out[75]:
```

	name	calories	protein	vitamins	rating
2	All-Bran	70	4	25	59.425505



Filtering Rows (1/3)

- We can also use the `[]` operator to apply conditions on one or more columns of the DataFrame.
- Rows of the DataFrame which satisfy those conditions are filtered out.

```
In [36]: condition = df['calories'] > 70  
df[condition]
```

Out[36]:

	name	calories	protein	vitamins	rating
1	100% Natural Bran	120	3	0	33.983679
4	Almond Delight	110	2	25	34.384843

```
In [37]: df[ df['calories'] > 70]
```

Out[37]:

	name	calories	protein	vitamins	rating
1	100% Natural Bran	120	3	0	33.983679
4	Almond Delight	110	2	25	34.384843



Filtering Rows (2/3)

and (&)

- We can also group conditions using the and operator.
- Symbol for and in pandas is `&`. It works the same way as `and` in Python.
- Note: Each condition should be in parentheses.

```
In [38]: df[ (df['calories'] > 70) & (df['protein'] < 4)]
```

```
Out[38]:
```

	name	calories	protein	vitamins	rating
1	100% Natural Bran	120	3	0	33.983679
4	Almond Delight	110	2	25	34.384843



Filtering Rows (3/3)

or (|)

- We can also group conditions using the or operator.
- Symbol for and in pandas is | It works the same way as or in Python.
- Note: Each condition should be in parentheses.

```
In [39]: df[ (df['calories'] > 70) | (df['protein'] > 3)]
```

```
Out[39]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843



loc (1/4)

Indexing

- loc is used to index/slice a group of rows and columns based on their labels.
- The first argument is the row label and the second argument is the column label.
- In the following example we index the first row and the first column.

```
In [79]: df
```

```
Out[79]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843

```
In [88]: df.loc[0, 'name']
```

```
Out[88]: '100% Bran'
```



loc (2/4)

Indexing

- If we pass a list of row and column labels, we get a DataFrame.
- In the following example, we index first row and first column, but we pass the labels as lists. We get a DataFrame.

```
In [79]: df
```

```
Out[79]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843

```
In [89]: df.loc[[0], ['name']]
```

```
Out[89]:
```

	name
0	100% Bran



loc (3/4)

Slicing

- We can also slice rows and/or columns using the loc method.
- Both the start and stop index of a slice with loc are inclusive.
- In the following example, we slice the first 5 rows and the first 3 columns of the DataFrame. The result is a DataFrame.

```
In [41]: df
```

```
Out[41]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
5	Apple Cinnamon Cheerios	110	2	25	29.509541
6	Apple Jacks	110	2	25	33.174094
7	Basic 4	130	3	25	37.038562
8	Bran Chex	90	2	25	49.120253
9	Bran Flakes	90	3	25	53.313813

```
In [42]: df.loc[0:4, 'name':'protein']
```

```
Out[42]:
```

	name	calories	protein
0	100% Bran	70	4
1	100% Natural Bran	120	3
2	All-Bran	70	4
3	All-Bran with Extra Fiber	50	4
4	Almond Delight	110	2



loc (4/4)

Indexing and Slicing

- We can index and slice simultaneously as well.
- In the following example we index rows and slice columns. The opposite is also possible.

```
In [43]: df.loc[[5, 8], 'name': 'protein']|
```

```
Out[43]:
```

	name	calories	protein
5	Apple Cinnamon Cheerios	110	2
8	Bran Chex	90	2



Quiz Time

1. Consider the given DataFrame called df. What will be the result of the following loc command?
`df.loc[[0, 1], ['name']]`
 - a) DataFrame
 - b) Series
 - c) Cell



Quiz Time

1. Consider the given DataFrame called df. What will be the result of the following loc command?
`df.loc[[0, 1], ['name']]`
 - a) DataFrame
 - b) Series
 - c) Cell



iloc (1/4)

Indexing

- iloc is used to index/slice a group of rows and columns.
- Iloc takes row and column positions as arguments and not their labels.
- The first argument is the row position and the second argument is the column position.
- In the following example we index the 10th row and the third column. The result is a Series.

```
In [41]: df
```

```
Out[41]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
5	Apple Cinnamon Cheerios	110	2	25	29.509541
6	Apple Jacks	110	2	25	33.174094
7	Basic 4	130	3	25	37.038562
8	Bran Chex	90	2	25	49.120253
9	Bran Flakes	90	3	25	53.313813

```
In [44]: df.iloc[9, 2]
```

```
Out[44]: 3
```




iloc (2/4)

Indexing

- If we pass a list of row and column positions, we get a DataFrame.
- In the following example, we index 10th row and third column, but we pass the positions as lists. We get a DataFrame.

```
In [45]: df.iloc[[9], [2]]
```

```
Out[45]:
```

protein

9

3



iloc (3/4)

Slicing

- We can also slice rows and/or columns using the iloc method.
- We provide row and column positions for slicing using iloc.
- The start index of a slice with iloc is inclusive. However, the end index is exclusive.
- In the following example, we slice the first 5 rows and the first 3 columns of the DataFrame. The result is a DataFrame.

```
In [46]: df.iloc[0:5, 0:3]
```

```
Out[46]:
```

	name	calories	protein
0	100% Bran	70	4
1	100% Natural Bran	120	3
2	All-Bran	70	4
3	All-Bran with Extra Fiber	50	4
4	Almond Delight	110	2



iloc (4/4)

Indexing and Slicing

- We can index and slice simultaneously as well.
- In the following example we index rows and slice columns. The opposite is also possible.

```
In [47]: df.iloc[[0, 2, 4], 0:3]
```

```
Out[47]:
```

	name	calories	protein
0	100% Bran	70	4
2	All-Bran	70	4
4	Almond Delight	110	2



Quiz Time

1. Consider the given DataFrame called df. What will be the result of the following iloc command?
`df.loc[[0, 2], [2]]`
 - a) DataFrame
 - b) Series
 - c) Cell
1. The stop index in iloc slice is inclusive. Is this statement True or False?
 - a) True
 - b) False



Quiz Time

1. Consider the given DataFrame called df. What will be the result of the following iloc command?
`df.loc[[0, 2], [2]]`
 - a) DataFrame
 - b) Series
 - c) Cell
1. The stop index in iloc slice is inclusive. Is this statement True or False?
 - a) True
 - b) False



Adding and Deleting Rows and Columns (1/4)

Adding Rows

- We can add more rows to our DataFrame using the loc method.
- If the row label does not exist, a new row with the specified label will be added at the end of the row.

In [24]: df

Out[24]:

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843

In [68]: df.loc[6] = ['Trix', 110, 1, 25, 27.753301]
df

Out[68]:

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
6	Trix	110	1	25	27.753301



Adding and Deleting Rows and Columns (2/4)

Deleting Rows

- We can delete rows from the DataFrame using drop() function by specifying axis=0 for rows.
- Provide the labels of the rows to be deleted as argument to the drop() function.
- Don't forget to use inplace=True, otherwise the original DataFrame will remain unchanged.

```
In [69]: df.drop(2, axis=0, inplace=True)
```

```
In [70]: df
```

```
Out[70]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
6	Trix	110	1	25	27.753301



Adding and Deleting Rows and Columns (3/4)

Adding Columns

- To add a column to the DataFrame, we use the same notation as adding a key, value pair to a dictionary.
- Instead of the key, we provide column name in the square brackets, and then provide a list of values for that column.
- If no column with the given name exists, a new column with the specified name and values will be added to the DataFrame.

```
In [71]: df['My Column'] = ['A', 'B', 'C', 'D', 'E']  
df
```

```
Out[71]:
```

	name	calories	protein	vitamins	rating	My Column
0	100% Bran	70	4	25	68.402973	A
1	100% Natural Bran	120	3	0	33.983679	B
3	All-Bran with Extra Fiber	50	4	25	93.704912	C
4	Almond Delight	110	2	25	34.384843	D
6	Trix	110	1	25	27.753301	E



Adding and Deleting Rows and Columns (4/4)

Deleting Columns

- We can also delete columns of the DataFrame using drop() function by specifying axis=1 for columns.
- Provide the column names to be deleted as argument to the drop() function.
- Don't forget to use inplace=True, otherwise the original DataFrame will remain unchanged.

```
In [72]: df.drop('My Column', axis=1, inplace=True)
```

```
In [73]: df
```

```
Out[73]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
6	Trix	110	1	25	27.753301



Sorting Values (1/2)

Ascending

- We can sort the values of a DataFrame with respect to a column using the `sort_values()` function, which sorts the values in ascending order by default.
- If the values of the column are alphabets, they are sorted alphabetically.
- If the values of the column are numbers, they are sorted numerically.

```
In [74]: df.sort_values(by='calories')
```

```
Out[74]:
```

	name	calories	protein	vitamins	rating
3	All-Bran with Extra Fiber	50	4	25	93.704912
0	100% Bran	70	4	25	68.402973
4	Almond Delight	110	2	25	34.384843
6	Trix	110	1	25	27.753301
1	100% Natural Bran	120	3	0	33.983679



Sorting Values (2/2)

Descending

- To sort the values in descending order, we set ascending = False in the sort_values() function.

```
In [75]: df.sort_values(by='calories', ascending=False)
```

```
Out[75]:
```

	name	calories	protein	vitamins	rating
1	100% Natural Bran	120	3	0	33.983679
4	Almond Delight	110	2	25	34.384843
6	Trix	110	1	25	27.753301
0	100% Bran	70	4	25	68.402973
3	All-Bran with Extra Fiber	50	4	25	93.704912



Exporting and Saving Pandas DataFrame

- To export a DataFrame as a csv file, use `to_csv()` function.
- If a file with the specified filename exists, it will be modified. Otherwise, a new file with the specified filename will be created.
- If you do not want to store index column in the csv file, you can set `index_label=False` in the `to_csv()` function.

```
In [76]: df.to_csv('myFile.csv', index_label=False)
```

```
In [77]: newDf = pd.read_csv('myFile.csv')
newDf
```

Out[77]:

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
6	Trix	110	1	25	27.753301



Concatenating DataFrames (1/3)

- We can concatenate two or more DataFrames together using `pandas.concat()` function.

First Data
Frame

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843

Second Data
Frame

	name	calories	protein	vitamins	rating
5	Apple Cinnamon Cheerios	110	2	25	29.509541
6	Apple Jacks	110	2	25	33.174094
7	Basic 4	130	3	25	37.038562
8	Bran Chex	90	2	25	49.120253
9	Bran Flakes	90	3	25	53.313813
10	Cap'n'Crunch	120	1	25	18.042851

Resultant Data Frame

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505
3	All-Bran with Extra Fiber	50	4	25	93.704912
4	Almond Delight	110	2	25	34.384843
5	Apple Cinnamon Cheerios	110	2	25	29.509541
6	Apple Jacks	110	2	25	33.174094
7	Basic 4	130	3	25	37.038562
8	Bran Chex	90	2	25	49.120253
9	Bran Flakes	90	3	25	53.313813
10	Cap'n'Crunch	120	1	25	18.042851



Concatenating DataFrames (2/3)

- We can also concatenate two or more DataFrames side-by-side each other.

First Data
Frame

	name	calories	protein	vitamins	rating
0	Apple Cinnamon Cheerios	110	2	25	29.509541
1	Apple Jacks	110	2	25	33.174094
2	Basic 4	130	3	25	37.038562

Second Data
Frame

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505

	name	calories	protein	vitamins	rating		name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973		Apple Cinnamon Cheerios	110	2	25	29.509541
1	100% Natural Bran	120	3	0	33.983679		Apple Jacks	110	2	25	33.174094
2	All-Bran	70	4	25	59.425505		Basic 4	130	3	25	37.038562

Resultant Data Frame



Concatenating DataFrames (3/3)

- To join two or more DataFrames side-by-side, use `axis = 1` in the `pandas.concat()` function.

```
In [26]: df
```

```
Out[26]:
```

	name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973
1	100% Natural Bran	120	3	0	33.983679
2	All-Bran	70	4	25	59.425505

```
In [36]: df2
```

```
Out[36]:
```

	name	calories	protein	vitamins	rating
0	Apple Cinnamon Cheerios	110	2	25	29.509541
1	Apple Jacks	110	2	25	33.174094
2	Basic 4	130	3	25	37.038562

```
In [37]: pd.concat([df, df2], axis=1)
```

```
Out[37]:
```

	name	calories	protein	vitamins	rating		name	calories	protein	vitamins	rating
0	100% Bran	70	4	25	68.402973		Apple Cinnamon Cheerios	110	2	25	29.509541
1	100% Natural Bran	120	3	0	33.983679		Apple Jacks	110	2	25	33.174094
2	All-Bran	70	4	25	59.425505		Basic 4	130	3	25	37.038562



groupby() (1/7)

- groupby() function is used to group DataFrame based on Series.
 - The DataFrame is splitted into groups.
 - An aggregate function is applied to each column of the splitted DataFrame.
 - Results are combined together.
- Consider the following DataFrame.

	Gender	Score
0	female	85
1	male	88
2	female	95
3	male	80



groupby() (2/7)

- The 'Gender' column contains two values, male and female.
- Let's split our DataFrame into two parts based on 'Gender' column;
 - First part will contain the rows where Gender = male
 - Second part will contain the rows where Gender = female

	Gender	Score
0	female	85
2	female	95

	Gender	Score
1	male	88
3	male	80



groupby() (3/7)

- If we find the mean score of both the genders, this is what we get.

Score	
Gender	
female	90

Score	
Gender	
male	84



groupby() (4/7)

- Let's combine the two results together. This is what we get.

Score	
Gender	
female	90
male	84



groupby() (5/7)

- The groupby() function works exactly the same way, except that it makes things easier for us.
- In the given example, we group our DataFrame on the basis of 'Gender' column, and then apply the aggregate function mean() on it.

```
In [55]: df
```

```
Out[55]:
```

	Gender	Score
0	female	85
1	male	88
2	female	95
3	male	80

```
In [56]: df.groupby(x['Gender']).mean()
```

```
Out[56]:
```

	Score
Gender	
female	90
male	84



groupby() (6/7)

- Note that aggregate functions are applied automatically on all the columns of the DataFrame except the one used to group the DataFrame.

In [72]:

```
df
```

Out[72]:

	Gender	Math	English
0	female	85	80
1	male	88	88
2	female	95	92
3	male	80	95

In [73]:

```
df.groupby(x['Gender']).mean()
```

Out[73]:

	Math	English
Gender		
female	90.0	86.0
male	84.0	91.5



groupby() (7/7)

- The common aggregate functions are;
 - `mean()`
 - `sum()`
 - `max()`
 - `min()`
 - `median()`
 - `count()`
 - `std()` (standard deviation)



Resources

- <https://www.geeksforgeeks.org/python-pandas-dataframe/?ref=lbp>
- https://pandas.pydata.org/pandas-docs/stable/user_guide/dsintro.html
- https://www.w3schools.com/python/pandas/pandas_dataframes.asp